
What the Final Patch Hides: Event-Level Regression Measurement for Coding-Agent Harnesses

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 Coding agents are evaluated on the patch they leave behind. We show that this
2 final state hides most of what happens during a run. We built an instrument that
3 records every mutating action of an agent as a git commit on a hidden reference,
4 replays each benchmark instance’s held-out tests at every recorded state inside
5 the benchmark’s own container image, and attributes each detected regression to
6 the harness check that could have caught it. Applied to SWE-bench Verified, the
7 instrument recorded 184 regression events across the runs that produced any; the
8 final patch exposed one. Regression frequency depended on the model stack: on the
9 same 40 instances under the same harness, one frontier stack produced no events
10 in 40 runs and another produced 140 events in 11 incidents (Fisher $p = 0.010$).
11 A pre-registered comparison of recovery policies, analysed with code committed
12 before unblinding, found that rollback to a verified checkpoint left contamination
13 in one final tree where no recovery left it in nine (paired sign test, $p = 0.022$),
14 but rollback also resolved fewer tasks, because the verifier that triggers it rejected
15 patches the official grader accepted. Replacing that verifier with the repository’s
16 own tests removed the tradeoff: 65% resolved with no contaminated final tree
17 (McNemar $p = 0.012$ against plain rollback), and a variant that reads only half of
18 those tests, graded on the other half, resolved 70% with no held-out failure. We
19 release the instrument, its validation protocol, and a catalogue of 28 measurement
20 defects encountered while building it, 23 of which produced plausible numbers
21 rather than errors.

22 1 Introduction

23 Benchmarks for coding agents grade the final patch. SWE-bench [1] and its Verified subset [2] run
24 the repository’s tests against the submitted patch and report whether the target tests now pass and the
25 previously passing tests still pass. Work on agent-introduced regressions follows the same convention
26 and counts the previously passing tests the submitted patch breaks [8].

27 This convention has a blind spot. A regression that an agent introduces and later repairs leaves no
28 trace in the final patch, and any harness with a recovery mechanism produces that pattern by design.
29 Whether it matters turns on two unanswered questions: how often regressions occur during a run, and
30 how much of that activity survives to the final state.

31 We answer both by measuring the timeline rather than the endpoint. Our contributions are:

- 32 1. **An instrument for event-level regression measurement.** Every mutating tool call is
33 recorded as a commit on a git reference the agent cannot see; after the run, the instance’s

held-out passing tests are replayed at every recorded state inside the benchmark’s pinned container, and detection is attributed by coverage. Validity is established by controls, a liveness gate, a golden check through the agent’s own execution path, and a re-scoring control (Section 3).

2. **Measurements on SWE-bench Verified.** Across the GPT-5.6 rollback runs that produced any regression, the timeline recorded 184 events and the final state exposed one (Section 5.2). Event frequency depended on the model stack: identical instances and harness gave 0 events under one frontier stack and 140 under another (Section 5.3).
3. **A pre-registered comparison of recovery policies, and a fix.** With the analysis committed before unblinding, rollback to a verified checkpoint left fewer contaminated final trees than no recovery ($p = 0.022$) but resolved fewer tasks; we trace the loss to verifier precision, and a post hoc regression-gated verifier removes it, reaching 65% resolve with no contaminated tree, 70% in a split variant graded on tests the gate never reads (Sections 5.4 to 5.6).
4. **A catalogue of 28 measurement defects** found while building the instrument, classified by mechanism (Appendix A); 23 produced a plausible number rather than an error.

All measurements are exploratory and were made on a development slice of the benchmark that is excluded from any confirmatory frame. No claim is made about the population of instances beyond that slice.

2 Background and related work

Benchmarks and their validity. SWE-bench [1] grades a patch by tests that must go from failing to passing and tests that must keep passing; Verified [2] is the human-screened 500-instance subset. Concerns about the static subset have accumulated: contamination, memorisation, and leaked solutions [6], flawed tests [7], and rolling [3, 5] or private [4] alternatives in response. UTBoost [7] showed the held-out tests are often insufficient, changing 24% of Verified leaderboard verdicts. TDAD [8] measured regressions in submitted patches on Verified and found them in about 30% of resolved instances; we measure the same quantity on the timeline, and Section 5.2 compares the two on the same runs.

Agent scaffolds and recovery. SWE-agent [9], OpenHands [10], Agentless [11], and mini-swe-agent [12] span the scaffold design space, built on CodeAct [13] and ReAct [14]. Within-episode recovery has been studied as self-reflection [15, 16], progress-gated recovery [17], checkpoint repair [18], and provenance-based rollback [19]; operating-system framings [20] treat checkpointing as a shared primitive. Process reward models and rubric verifiers [24, 25] approach the verifier-precision problem of Section 5.5 from the training side.

Automated program repair. The regression problem is the plausible-versus-correct patch problem of program repair [26, 28, 29]; regression tests are the main defence against overfitting patches [27]. Our results concern regressions the agent’s own process repairs before the patch is tested; trajectory-level evaluation [21–23] argues that resolve rate alone is uninformative. Public leaderboard trajectories record actions and observations but not the tree at each step, so this measurement cannot be mined from them without re-execution; the instrument commits the tree.

3 Instrument

Figure 1 shows the harness and the instrument. **Observational timeline.** After every mutating tool call, the working tree is committed to a git reference outside the agent’s view, using a private index so that the agent’s own `git status` and `git diff` are unchanged. Rollbacks and the end of the run are observations too, and because observations are commits, an archived run can be re-measured later without re-running the agent.

Exhaustive replay. For each observation, the instance’s passing-test set is run inside the pinned container against that observation’s tree; a regression event is a test that passes at one observation and fails at a later one. Every observation is replayed rather than bisected, because bisection assumes monotone verdicts and a recovery policy violates that by construction. Replay is scoped to the files holding the held-out tests, which is what keeps it affordable: 4 to 8 seconds per observation on the




Figure 1: The harness and the instrument. The planner, worker, monitor, and recovery policy run inside the benchmark’s pinned container; every mutating tool call is committed to a hidden git reference; held-out tests are replayed at every observation and detections are attributed by coverage. In the gated arm (dashed), the timeline’s tests serve as the monitor. The official grader sees only the final patch.

two repository families we timed (pytest and django), about 26 CPU-minutes for a 40-instance sweep, and a projected 80 CPU-hours for 500 instances at 100 edits each. Infrastructure failures during replay are recorded as missing observations rather than as test failures, and baseline-dead tests are excluded from event counting.

Attribution. A detected regression is classified three ways: *attributed* if a failing harness check and the broken held-out test both exercise a file the agent changed at that observation (using a per-instance coverage map built at the base commit), *co-occurring* if some harness check failed while the regression was open but no such link exists (an over-count of detection), and *unknown* if the coverage map cannot say.

Execution environment. The agent’s tools and the harness’s checks execute inside the same pinned container as the replay, with file changes synchronised to the host tree the timeline records (Appendix A.2).

Validation. Five gates run before any paid experiment: a negative control, a positive control with injected regressions including recovered ones, a flake screen, an unknown-rate ceiling, and a baseline liveness check. A golden check drives the gold patch through the agent’s real tool path and requires the grader to return "resolved" (a null run "unresolved"); it passed on three repository families. A re-scoring control re-measures an archived timeline under a changed instrument with no model calls; on both runs to date it reproduced the archived episode counts.

4 Experimental setup

Benchmark and slice. SWE-bench Verified, 500 instances. A 40-instance development slice (16 from django), stratified and fixed before any measurement, was used throughout and is excluded from any confirmatory frame. The GPT-5.6 stack also ran on the slice’s first 10 instances as a calibration, and plain rollback ran twice.

Harness. A planner decomposes the task into steps with verification commands; a worker executes each step with three tools (read, write, shell); a monitor runs the verification, and on failure the recovery policy acts: **rollback** (reset to the last verified checkpoint, retry with feedback), **repair-in-place** (retry from the failed tree), or **no recovery** (keep the failed step’s tree). Each run has a \$4 work-cost cap; exhaustion is scored as failure.

Models. Two frontier stacks under identical harness, policy, instances, and caps, named by their API model strings: claude-opus-4-7 (planner) with claude-sonnet-4-6 (worker), and




Figure 2: Every run that produced at least one regression event, in both GPT-5.6 sweeps. Left bars (red) are events recorded on the timeline; right bars (grey) are held-out test failures visible in the graded final patch.

114 gpt-5.6-sol (planner) with gpt-5.6-terra (worker); identifiers are recorded in every run mani-
115 fest.

116 **Outcomes.** Resolve is the official grader’s verdict on the final patch. Events are reported at three
117 levels because the distribution is overdispersed: distinct incidents (observations at which at least one
118 test broke), declared events (one per test function and onset, parametrised variants collapsed), and
119 bearing runs. Final-state contamination is the number of held-out tests failing in the graded final
120 patch, net of baseline-dead tests.

121 **Pre-declaration.** The event unit, the gates, and the contrast were declared before the relevant data
122 existed; the contrast’s analysis (exact paired sign tests, final-state contamination primary, incident
123 exposure co-primary) was committed before either comparison arm finished. One amendment,
124 making the gates conditional on the model stack, preceded unblinding (Appendix B).

125 5 Results

126 5.1 Resolve and cost

127 Under rollback, the Claude stack resolved 13 of 40 instances (32.5%, exact 95% CI 18.6% to 49.1%)
128 for a total sweep cost of \$34.24; the GPT-5.6 stack resolved 13 of 40 (32.5%; 13 of the 35 graded,
129 37.1%) for \$8.14. Three GPT-5.6 cells were never attempted (circuit breaker) and two were lost to a
130 file-synchronisation defect; all five count as unresolved out of 40 and are excluded from the graded
131 denominator. One instance cannot be resolved by any arm under the current official parser (Appendix
132 A.3), so every rate has a ceiling of 39. Runs are short: the median run makes 5 mutating tool calls
133 (IQR 4 to 8, maximum 19 over 235 runs; 94% make 10 or fewer), an order of magnitude below
134 public-scaffold trajectories. Paired comparisons use only instances graded in both arms.

135 5.2 The final state hides almost all regression activity

136 Figure 2 shows every run that produced a regression event under the GPT-5.6 stack with the rollback
137 policy, pooling the 10-instance calibration and the 40-instance sweep (47 runs). The timeline recorded
138 184 declared events across eight runs; the final patch exposed one. Run-weighted, seven of the eight
139 bearing runs ended with a clean final patch; three storms (73, 48, and 41 events) supply 88% of the
140 event count, so the run-weighted figure is the more robust statement.

141 Figure 3 shows a single run in detail. The grader marked it resolved with all graded tests passing; its
142 timeline contains three regressions of the same test function at observations 3, 5, and 7, each repaired
143 by rollback. Two of the three went undetected by any harness check while open. Across the full
144 GPT-5.6 sweep, 48 of 162 raw episodes had no co-occurring harness failure at all.

145 5.3 Regression frequency depends on the model stack, not the benchmark

146 On the same 40 instances, under the same harness and rollback policy, the Claude stack produced
147 no regression events in 40 runs (227 observations, all replayed). The GPT-5.6 stack produced 140
148 declared events in 11 incidents across 6 of 37 attempted runs (bearing fraction 16.2%, exact CI 6.2%
149 to 32.0%). A one-sided Fisher exact test on the bearing fractions gives $p = 0.010$; this is a single




Figure 3: One run that the official grader marked resolved, with all 147 graded tests passing. The timeline shows three regressions of the same test function, each repaired by rollback two observations later.




Figure 4: The five recovery arms under the GPT-5.6 stack on the same 40 instances. Contamination counts cells whose final patch fails a previously-passing test.

pairwise test on six bearing runs (reassign three of them and $p = 0.11$), and the two stacks also traverse different provider adapters. Observation density differs by a factor of 1.2, which does not account for 140 to 0. The two stacks resolved the same number of instances (13 of 40 each) for total sweep costs of \$34.24 against \$8.14: the stack that never broke adjacent behaviour cost four times as much per run. The three storms that supply 88% of the events are model edits, not harness artifacts: a six-line change to seaborn’s `plot.py`, a nine-line change to django’s `query.py`, and a 262-line deletion in flask’s `blueprints.py`, none truncated, none containing a shell idiom, and none from a capped turn (defect D10 was fixed before these sweeps).

The Claude result is a measured zero rather than a detection failure: the same stack’s earlier canary run produced three events, and re-scoring its archived timeline under the instrument used for the 40-run sweep reproduces them.

We had first read the low event rate as a property of the benchmark; the second stack falsified that. The pre-declared event-rate gate still failed for both stacks, so no arm here is a confirmatory pass (Appendix B), and we claim no more than a single-sweep dependence on the model stack.

5.4 Recovery policy: rollback keeps the final tree clean

Figure 4 summarises the pre-declared contrast. The primary endpoint, final-state contamination, is paired by instance: no recovery was worse than rollback on 9 instances, better on 1, and tied on 25 (exact sign test, $p = 0.022$). Repair-in-place was worse on 5, better on 1, and tied on 29 ($p = 0.22$). The co-primary endpoint, incident exposure, did not differ detectably for either comparison (no recovery $p = 0.45$; repair-in-place $p = 1.0$): regressions occurred at similar rates under every policy, and the policy determined whether they persisted.

Disclosure. The first unblinding gave $p = 0.375$. Seven cells (five no-recovery, two repair-in-place) whose final trees made the suite uncollectable had been dropped as ungradable, removing the most contaminated states from the endpoint they bore on most, whereas official grading scores such a patch as failing every test. The grader was corrected to distinguish an environment failure from a patch that kills the suite, and the seven archived trees were re-graded without re-running any agent, giving $p = 0.022$. The rule changed after the data were seen (Appendix B).

5.5 Rollback loses resolution to verifier precision

Rollback resolved fewer tasks than either alternative (Figure 4, left): 13 of 40 against 24 of 40 for repair-in-place and 22 of 40 for no recovery; paired by instance, it won one discordant pair and lost nine against repair-in-place (McNemar exact $p = 0.022$) and won three and lost nine against no recovery ($p = 0.15$). The loss is attributable to the verifier: of the 18 rollback runs that failed, 15 failed at the first step, and 9 of the 18 were resolved under a policy that kept the rejected work; the monitor’s planner-written check had rejected a patch the grader accepted, and rollback discarded it.

5.6 Regression-gated rollback removes the tradeoff

A fourth arm, added after the contrast was unblinded and therefore exploratory, replaces the monitor’s planner-written check with the repository’s previously-passing tests, run in the agent’s container at the start and after every attempt; a step is rejected only if a test that passed at the start fails now. It resolved **26 of 40 instances (65.0%)**, the most of any arm, with **no contaminated final tree** (Figure 4). Paired against plain rollback on the 35 shared instances, it resolved 10 that plain rollback did not and lost 1 (McNemar exact $p = 0.012$; post hoc, one of eight tests reported, unadjusted; paired onset exposure $p = 0.73$). The gate reads the instance’s previously-passing tests, selected from benchmark metadata, so its baseline oracle coincides with the grading oracle (median ratio 1.00) and the clean final trees are guaranteed by construction. Both gated arms also *select* their tests from that metadata, which observes roughly the gold test patch’s blast radius, information a deployed agent lacks; the deployed analogue gates on the full repository suite at baseline (django’s full suite takes about two minutes per attempt at four workers, against 4 to 8 seconds scoped, with 97 baseline failures and errors in the pinned image that a flake screen must absorb), and the gated resolve rates are upper bounds on a metadata-free gate. A fifth arm, a split variant, removes that guarantee: the gate reads results for a deterministic half of the previously-passing test ids (2,278) and never reads the other half (2,585), on which the grader then rules. It resolved **28 of 40 (70.0%)**, indistinguishable from the full gate (5 discordant pairs against 3, $p = 0.73$) and 13 pairs against 2 over plain rollback (McNemar exact $p = 0.007$), with no final tree failing a held-out test (the one apparent exception is the parser-capped instance of Section 5.1). A second plain-rollback sweep reproduced the first (13 and 13 resolved; 6 of 37 bearing both times), so the gain is outside run-to-run variability. Running the tests before accepting an edit is standard advice; the contribution is measuring what it buys and why plain rollback loses it.

6 Discussion

Why timeline regressions matter. A regression the agent repairs before submission did not reach the user, but it is not free: across the bearing runs, 27% of spend (\$0.63 of the \$2.33 spent across the eight bearing runs) went to work done while a regression was open. The events are the per-step ground truth that process reward models for software agents [24, 25] need; for agent-OS design, a recovery primitive without an observation primitive hides exactly the activity that distinguishes recovery policies.

Summary. Final-state evaluation misses the regressions an agent’s own process repairs, nearly all of them in our measurements; measuring the timeline is feasible, validates against controls, separates recovery policies the final patch cannot, and locates rollback’s cost in verifier precision. Building the instrument produced 28 defects (Appendix A); two rules, record infrastructure failures as missing observations and require a positive liveness signal, would have prevented most of them.

7 Limitations

All measurements come from a 40-instance development slice, one sweep per arm, with variability measured only for plain rollback under GPT-5.6. The cross-stack comparison is six bearing runs, confounded with the provider adapter; the gated arms are post hoc and select their tests from benchmark metadata. Runs are an order of magnitude shorter than public-scaffold trajectories (median 5 mutating calls against tens to hundreds), so recovery dynamics at length are unmeasured. The held-out tests observe roughly the gold test patch’s blast radius, so event counts are lower bounds; intervals ignore repository clustering; and the instrument has not been applied to a public scaffold, which is the next experiment.

References

- [1] C. Jimenez et al. SWE-bench: Can language models resolve real-world GitHub issues? ICLR 2024. arXiv:2310.06770.
- [2] OpenAI. Introducing SWE-bench Verified. 2024.
- [3] L. Zhang et al. SWE-bench Goes Live! arXiv:2505.23419, 2025.
- [4] X. Deng et al. SWE-Bench Pro: Can AI agents solve long-horizon software engineering tasks? arXiv:2509.16941, 2025.
- [5] I. Badertdinov et al. SWE-rebench: An automated pipeline for task collection and decontaminated evaluation of software engineering agents. arXiv:2505.20411, 2025.
- [6] S. Liang et al. The SWE-Bench Illusion: When state-of-the-art LLMs remember instead of reason. arXiv:2506.12286, 2025.
- [7] H. Yu et al. UTBoost: Rigorous evaluation of coding agents on SWE-Bench. ACL 2025. arXiv:2506.09289.
- [8] TDAD: Regressions introduced by coding agents on SWE-bench Verified. arXiv:2603.17973, 2026.
- [9] J. Yang et al. SWE-agent: Agent-computer interfaces enable automated software engineering. NeurIPS 2024. arXiv:2405.15793.
- [10] X. Wang et al. OpenHands: An open platform for AI software developers as generalist agents. ICLR 2025. arXiv:2407.16741.
- [11] C. Xia et al. Agentless: Demystifying LLM-based software engineering agents. arXiv:2407.01489, 2024.
- [12] SWE-agent team. mini-swe-agent. Software, 2025.
- [13] X. Wang et al. Executable code actions elicit better LLM agents. ICML 2024. arXiv:2402.01030.
- [14] S. Yao et al. ReAct: Synergizing reasoning and acting in language models. ICLR 2023. arXiv:2210.03629.
- [15] N. Shinn et al. Reflexion: Language agents with verbal reinforcement learning. NeurIPS 2023. arXiv:2303.11366.
- [16] A. Madaan et al. Self-Refine: Iterative refinement with self-feedback. NeurIPS 2023. arXiv:2303.17651.
- [17] ReflexGrad: Within-episode failure recovery in LLM agents via progress-gated dual-process routing. arXiv:2511.14584, 2025.
- [18] REPOT: Recoverable program-of-thought via checkpoint repair. arXiv:2605.30052, 2026.
- [19] From agent traces to trust: Evidence tracing and execution provenance in LLM agents. arXiv:2606.04990, 2026.
- [20] K. Mei et al. AIOS: LLM agent operating system. arXiv:2403.16971, 2024.
- [21] Holistic Agent Leaderboard: The missing infrastructure for AI agent evaluation. arXiv:2510.11977, 2025.
- [22] What resolve rate hides: Trajectory structure diagnostics for coding agents. arXiv:2607.06184, 2026.
- [23] SWE-EVO: Benchmarking coding agents in long-horizon software evolution scenarios. arXiv:2512.18470, 2025.
- [24] Agentic rubrics as contextual verifiers for SWE agents. arXiv:2601.04171, 2026.

- 270 [25] SWE-Shepherd: Advancing PRMs for reinforcing code agents. arXiv:2604.10493, 2026.
- 271 [26] Patch overfitting in program repair: A survey. 2024.
- 272 [27] Y. Wang et al. When automated program repair meets regression testing: An extensive study on
273 two million patches. ACM TOSEM, 2024. arXiv:2105.07311.
- 274 [28] Patch correctness assessment: A survey. ACM TOSEM, 2024. doi:10.1145/3702972.
- 275 [29] H. Ye et al. Automated patch assessment for program repair at scale. Empirical Software
276 Engineering, 2021.
- 277 [30] SWE-bench issue #601: Test result hijacking via stdout forging in the evaluation harness.
278 github.com/SWE-bench/SWE-bench, 2025.
- 279 [31] OpenHands issues #4235 and #7044: Environment errors in SWE-bench instance evaluation.
280 github.com/OpenHands/OpenHands.
- 281 [32] J. Yang et al. SWE-smith: Scaling data for software engineering agents. arXiv:2504.21798,
282 2025.
- 283 [33] J. Pan et al. Training software engineering agents and verifiers with SWE-Gym.
284 arXiv:2412.21139, 2024.
- 285 [34] Context as a tool: Context management for long-horizon SWE agents. arXiv:2512.22087, 2025.

286 **A The measurement defect catalogue**

287 Two rules would have prevented most of the defects below: record infrastructure failures as missing
288 observations, and require a positive liveness signal.

289 **A.1 The catalogue by mechanism**

290 Each row is a defect found while building the instrument. **S**: it produced a plausible number rather
291 than an error. **G**: it was present while the test suite passed (A4 is the suite). **H**: it was findable only
292 on a clean host. Class A rows depend on the machine the code runs on; class B rows render an
293 infrastructure failure as a measurement; class C rows measure a state where an event was defined;
294 class D rows consume a producer's output without checking the producer.

#	Defect	Consequence	S	G	H
A1	Shadow commits inherited the machine’s git identity	timeline silently empty on any clean machine	✓	✓	✓
A2	Gate probe invoked bare <code>python</code>	every probe exit 127 on a clean host	✗	✓	✓
A3	Unpinned SDK resolved a different major version	two machines ran different code	✗	✓	✓
A4	Test fixtures verified with bare <code>pytest</code> from PATH	15 tests fail on a clean host as harness defects	✗	—	✓
A5	A benchmark scorer invoked bare <code>python</code>	score 0.0 from a missing interpreter	✓	✓	✓
A6	Agent executed on the host; measurement in the container	26 of 28 zero-step runs; see A.2	✓	✓	✓
B1	Output marker used shell :	a 13/13 run scored as 13 errors	✓	✓	
B2	Results on stderr, markers on stdout	one framework’s results outside the parsed slice	✓	✓	
B3	Probe diffed against a deleted commit	every graded test a hole	✓	✓	
B4	Interrupted mirror clone accepted with zero commits	later instances fail far from the cause	✓	✓	
B5	Container workdir unvalidated	every exec exit 127	✓	✓	
B6	Isolation removed loopback, not only the internet	23.5% of the oracle dead	✓	✓	
295 B7	Provider cache served removed containers	later cells replay against nothing and score clean	✓	✓	
B8	Tree reset reverted image build-time edits	one repository family’s oracle dead	✓	✓	
C1	Bisection assumed monotone verdicts	recovered regressions invisible	✓	✓	
C2	Declared event unit not implemented	rate off by up to $2.7\times$	✓	✓	
C3	Rollbacks produced no observation	recoveries invisible to the timeline	✓	✓	
C4	"Persists past the step boundary" undefined	one reading excludes the events of interest	✓	✓	
D1	Audit field never populated	0 tool errors reported, always	✓	✓	
D2	Malformed planner output crashed the run	33% of runs lost as task failures	✗	✓	
D3	Dependency install counted as agent work	attribution vacuous	✓	✓	
D4	Join took the first observation, not the last	failures pinned to the wrong tree	✓	✓	
D5	Executed config rebuilt from a name	manifest described a different run	✓	✓	
D6	Git handles never released	sweeps die after ~400 cells	✓	✓	
D7	Console re-walked the tree per run	presented as a dead server	✗	✓	
D8	Detection events lacked the ids attribution joined on	attributed detection always unknown	✓	✓	
D9	Absent coverage rendered as measured silence	unmeasured episodes reported as silent	✓	✓	
D10	Output-capped turns executed	half-written file; a 78-event storm	✓	✓	

296 Totals: 23 of 28 silent; 27 of 28 present under a passing suite; 6 of 28 findable only on a clean host.
297 Seven further defects found after this census closed are described where they arose (Section 5.4,
298 Appendix A.3, and the repository log).

299 A.2 The defect that invalidated a conclusion

300 After nineteen fixes, every validation gate passed, the re-scoring control reproduced archived episode
301 counts exactly, and three pilots (80 runs) measured an event rate far below the pre-declared gate.
302 We drafted the conclusion the rule directed: the benchmark offered too little opportunity, so switch
303 benchmarks. The conclusion was wrong. The harness ran the agent on the host in a bare source
304 checkout, while every probe and every gate ran inside the pinned container. On the host, `import`
305 `matplotlib` in that checkout succeeds by importing the uncompiled source tree as a namespace
306 package, and everything downstream fails in ways indistinguishable from an incompetent agent.
307 Twenty-six of 28 zero-step runs traced to this. The gates had validated the measurement path and
308 never the agent’s execution path. The fix was to route the agent’s tools and checks through the same
309 container as the replay, and to add a per-cell environment parity check that runs before any model
310 call.

311 A.3 Mechanisms in public infrastructure

312 Class A: OpenHands issues #4235 and #7044 [31], environment construction failures against SWE-
313 bench images. Class B: UTBoost’s finding that the held-out oracle is insufficient at leaderboard scale
314 [7], and, on SWE-bench Live, baseline tests that fail in the unmodified image because the calendar
315 passed a deprecation date encoded in the instance. Class C: final-state regression measurement [8].
316 Class D: the SWE-bench grader accepting forged test output on stdout [30], and, at the harness’s
317 current revision, a parser that keys parametrised ids by their full text while the dataset stores them
318 truncated at the first space (the earlier parser’s behaviour), so 16 previously-passing ids of one instance
319 in our slice match no output and every submission to it grades as unresolved.

B Pre-declaration and analysis

The event unit is one (test function, onset observation) pair with parametrised variants collapsed. The gates for proceeding on a substrate were an event rate of at least 0.30 per run and a bearing fraction of at least 25%. Both stacks failed the conjunction (bearing 0% and 16.2%); the gates were then re-declared per (substrate, model stack) before the contrast was unblinded, and no arm reported here is confirmatory. The contrast's primary endpoint is final-state contamination, paired by instance, exact two-sided sign test; the co-primary is incident exposure. The analysis script was committed before either comparison arm finished. The regression-gated arm was added after the contrast was unblinded and is exploratory. No multiplicity adjustment was planned or applied. Budget exhaustion is scored as failure. Instances whose baseline oracle records failures in the unmodified image are excluded from resolve-rate comparability claims, and their dead tests are typed as missing observations for event counting.

Grading correction between unblindings. The first unblinding of the recovery-policy contrast gave $p = 0.375$ for the primary endpoint: seven cells (five no-recovery, two repair-in-place) had been dropped as ungradable because their final trees made the suite uncollectable. Official grading scores such a patch as failing every test, so the most contaminated states had been dropped from the endpoint they bore on most. The grader now distinguishes an environment failure from a patch that kills the suite, and the seven archived trees were re-graded without re-running any agent.